

University of Waterloo

CS240 Spring 2016

Assignment 3

Due Date: Wednesday, June 15, at 5:00pm

Please read <http://www.student.cs.uwaterloo.ca/~cs240/s16/guidelines.pdf> for guidelines on submission. This assignment contains only written problems. Submit your written solutions electronically as a PDF with file name a03wp.pdf using MarkUs. We will also accept individual question files named a03q1w.pdf, a03q2w.pdf, ... , a03q4w.pdf if you wish to submit questions as you complete them.

Problem 1 [6 marks]

We define the relation $A < B$ to mean that all of the keys in AVL tree A are smaller than all the keys in AVL tree B .

Design an algorithm to concatenate/merge any two AVL trees for which $A < B$. The result should be a single valid AVL tree. The worst-case running time of your algorithm should be $O(h)$ where h is the maximal height of the two trees. Explain how your algorithm satisfies these requirements.

Problem 2 [3+3+3+5=14 marks]

Suppose we have BST T . For any node $u \in T$, we define the left and right subtrees as $left(u)$ and $right(u)$. We also define $H(u)$ to be the height of u , and $N(u)$ to be the number of nodes in the subtree rooted at u ($N(u) = 1$ if u is a leaf and $N(u) = 0$ if u is the empty subtree).

a) [3 + 3 + 3 marks] Which of the following trees must be of height $O(\log n)$? Prove or disprove.

1. There is a constant $c > 0$ such that for all nodes $u \in T$
 $H(left(u)) \leq H(right(u)) + c$.
2. There is a constant $c > 0$ such that for all nodes $u \in T$
 $H(left(u)) - c \leq H(right(u)) \leq H(left(u)) + c$.
3. Every internal node $u \in T$ has exactly two children.

b) [5 marks] Assume that there exists a constant $0 < c \leq 1$ such that for every node $u \in T$:

- $N(\text{left}(u)) \geq c \times N(\text{right}(u)) - 1$, and
- $N(\text{right}(u)) \geq c \times N(\text{left}(u)) - 1$.

Show that the tree has height $O(\log n)$. (Hint: The constant hidden in big-O should depend on c .)

Problem 3 [4+4=8 marks]

We wish to improve upon binary search for the sorted array A with n distinct values. Instead of just blindly selecting the median index $i = \frac{l+r}{2}$ from the range of indices $A[l, r]$, *interpolation search* guesses the index of key k by estimating “how far away it should be from l ”. The guess is calculated by interpolating between the left (l) and right (r) boundary values.

```
InterpSearch( $A[l, r], k$ )
 $A$ : an array
 $l$ : index of the left boundary
 $r$ : index of the right boundary
 $k$ : key to search for
1.  if  $A[l] > k \ || \ A[r] < k$ 
2.      return false
3.   $i \leftarrow l + \lfloor (r - l) \frac{k - A[l]}{A[r] - A[l]} \rfloor$ 
4.  if  $A[i] = k$ 
5.      return true
6.  else if  $A[i] < k$ 
7.      return InterpSearch( $A[i + 1, r], k$ )
8.  else
9.      return InterpSearch( $A[l, i - 1], k$ )
```

For example, let $A = (51, 58, 81, 87, 99)$ and $k = 87$. Starting the search in the range $l = 0$ and $r = 4$, then $i = 3$, and $A[i] = 87 = k$ so k is in A .

- a) [4 marks] Which array with n distinct values results in the best case search time? Be as general as possible. Show that search always terminates in $O(1)$ time for this array, regardless of whether the key is stored in the array or not.
- b) [4 marks] Which array values and search key yield the worst case search time? Give an array A with n values (i.e., by defining $A[i] = f(i)$ for some function f) that demonstrates the worst case search.

Problem 4 [5+6+4+4+5+(4)+(4)=24 marks]

Consider a procedure called *skipify* which consumes an array of n items in ascending order, and then deterministically, produces a skip list with $\log(n)$ levels: starting with the leftmost element $-\infty$, promote every other element from the previous level to the level above; as an exception, always promote $+\infty$ along with $-\infty$ to keep them the same height.

- [5 marks] Write pseudocode for *skipify* following the conventions and notation from the notes (Module 5, Slide 9/16). If you're latex-ing your solution, consider using an *algorithm* object (see *LBSLInsert* in A3.tex for an example). Your code should run in $\Theta(n)$ time—explain how this requirement is satisfied.
- [6 marks] Define a *balanced skip list* (BSL) as a skip list built by *skipify* but with exactly $n = 2^k - 1$ elements for some $k > 0$. Derive the best and worst case search times for a BSL (with $n = 2^k - 1$ elements). Assume that BSL search works exactly the same way as the randomly-generated skip lists from the lectures.

Suppose that we want to preserve the balanced structure of the BSL to guarantee the search times we derived in **part b**. This prevents us from modifying the BSL after building it with *skipify*, so there's no direct way to insert new elements.

Problem: design a data structure that provides search complexity similar to the BSL, but also allows insertion in a reasonable amount of time. Solution: since we can create as many BSLs as we want, we define a List of BSLs (LBSL) as follows.

An LBSL contains a set of BSLs, each containing $2^i - 1$ elements for some i values. For example, if $n = 11$, the LBSL contains a BSL with $2^3 - 1 = 7$ nodes, a BSL with $2^2 - 1 = 3$ nodes, and a BSL with $2^1 - 1 = 1$ nodes (Figure 1).




Figure 1: The LBSL produced after inserting elements (5, 17, 9, 12, 3, 1, 15, 41, 23, 4, 37).

To insert new elements, we merge smaller BSLs to make bigger ones:

LBSLInsert(e, L)

e : item to be inserted

L : an LBSL with $2^{k-1} < n \leq 2^k - 1$ elements

1. **for** $i \leftarrow k$ **to** 1 **do**
2. **if** L contains two BSLs A_1, A_2 of size $2^i - 1$
3. using *skipify*, replace A_1, A_2 with a BSL of size $2^{i+1} - 1$ containing A_1, A_2 and e
4. **return**
5. create a new BSL E of size 1 that contains only e
6. add E to L

- c) [4 marks] Create an empty LBSL and insert values 1 through 7 using *LBSLInsert*. Draw the LBSL after every insertion.
- d) [4 marks] Consider the worst case insertion time for an LBSL with n elements (not necessarily $n = 2^k - 1$). Describe the structure of the LBSL just before the insertion, and then derive the worst case insertion time.
- e) [5 marks] Give pseudocode for searching the LBSL and derive the worst case run-time. Also, discuss whether there's an advantage to searching the BSLs from largest to smallest, or vice versa.
- f) [Bonus: 4 marks] Insert calls *skipify* every time it replaces two BSLs size $2^i - 1$ with a new BSL size $2^{i+1} - 1$. Consider the total cost of all *skipify* calls after inserting values 1 through $n = 2^k - 1$. Show that this total cost is $\Theta(n \log n)$.
- g) [Bonus: 4 marks] Show that the LBSL insertion cost, amortized over n elements, is $\Theta(\log n)$. Note: we probably won't cover amortized analysis before this assignment is due. This bonus question is intended as an opportunity to do some independent investigation; in that regard, please do not ask questions about the bonus problems on Piazza.